

# Ankur Dave

---

2703 Ridge Rd Apt. 203 Berkeley, CA 94709  
[ankurdave@gmail.com](mailto:ankurdave@gmail.com) <http://ankurdave.com>  
mobile: 650-701-7705

## RESEARCH EXPERIENCE

---

### UC Berkeley AMPLab

Berkeley, CA

Undergraduate Research Assistant

September 2010 to May 2013

Contributed to Spark, an open source cluster computing framework written in Scala. Advised by Matei Zaharia and Scott Shenker.

- Led the design and development of Arthur, a distributed replay debugger for Spark programs; gave a talk at the AMPLab Winter 2012 retreat; and wrote a technical report [3]
- Used Spark to build Bagel, a Pregel-like graph processing framework, and presented a poster at the AMPLab Summer 2011 retreat
- Coauthored an NSDI 2012 paper on Spark [1] that won Best Paper Award, contributing Bagel and an evaluation of Spark's user-controllable partitioning

### Microsoft Research

Redmond, WA

Research Intern

June to August 2010

Designed and implemented CloudClustering, a scalable clustering algorithm on the Windows Azure cloud, for Microsoft Research's eXtreme Computing Group using C#. Generalized the design into a set of architectural patterns for data processing using cloud services. Published a workshop paper [2] and gave a talk at DataCloud 2011.

## INDUSTRY EXPERIENCE

---

### Google

Mountain View, CA

Software Engineer Intern

May to August 2012

Contributed to Google's workflow execution system, which handles collections of processes with dependencies.

- Designed and implemented an algorithm to improve efficiency at allocating cluster resources for complex workflows
- Productionized a tool for static analysis of user workflows to help design and prioritize new features
- Contributed cache-related bugfixes and features to Google's next-generation data processing substrate

### Facebook

Palo Alto, CA

Software Engineer Intern

May to August 2011

Added network usage tracking and limit enforcement to Facebook's cluster manager using C++ and Linux cgroups.

## PUBLICATIONS

---

### Conference and Workshop Papers

- [1] Matei Zaharia, Mosharaf Chowdhury, Tathagata Das, Ankur Dave, Justin Ma, Murphy McCauley, Michael Franklin, Scott Shenker, and Ion Stoica. [Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing](#), NSDI 2012, April 2012. *Best Paper Award and Honorable Mention for Community Award*.
- [2] Ankur Dave, Wei Lu, Jared Jackson, and Roger Barga. [CloudClustering: Toward an Iterative Data Processing Pattern on the Cloud](#). DataCloud 2011, May 2011.

### Technical Reports

- [3] Ankur Dave, Matei Zaharia, Scott Shenker, Ion Stoica. [Arthur: Rich Post-Facto Debugging for Production Analytics Applications](#). Submitted to USENIX ATC 2013, January 2013.
- [4] Ankur Dave. [Optimizing Boggle Boards: An Evaluation of Parallelizable Techniques](#). IB Extended Essay, January 2009.

## EDUCATION

---

### University of California, Berkeley

Starting August 2013

Ph.D. candidate, Computer Science

### University of California, Berkeley

August 2010 to May 2013

B.S., Electrical Engineering and Computer Science

GPA: 3.79/4.0

## SKILLS

---

- Programming languages: Scala, Python, C, C++, Scheme, C#, Java, PHP, JavaScript